
Music classification mini project: Instructions

January 24, 2017, Version 1.0

Andreas Svensson, andreas.svensson@it.uu.se

Department of Information Technology, Uppsala University

Abstract

This document contains the instructions for the classification mini project. The problem is to classify a set of 100 songs, and predict whether Andreas Svensson would like them or not, based on a training data set of 400 songs. You are expected to (i) try some (or all) classification methods that have been presented in the course and evaluate their performance on the problem, and (ii) make a decision which one to use and ‘put in production’ by uploading your predictions to a website, where your prediction will be evaluated and also compared to the performances of the other student groups. You will document your project by writing a report, which will be reviewed anonymously by your colleagues.

0 Requirements

The project is to be done in groups of 3-4 students. All tasks described in this document have to be done, in order to pass the project. To encourage discussions the project is carried out in groups, and of course *all group members have to take part in the project*.

1 Problem

The problem in this project is to tell which songs, in a dataset of 100 songs, Andreas Svensson is going to like (see Figure 1). The data set consists not of the songs themselves, but of high-level features (extracted using the web-API from Spotify¹). These high-level features describe characteristics such as the acousticness, danceability, energy, instrumentalness, valence and tempo of each song.



Figure 1: Andreas Svensson listening to music.

To your help, you are provided a training data set with 400 songs, each of which Andreas has labeled with LIKE or DISLIKE. You are expected to use all the knowledge that you have acquired in the course (so far) about classification algorithms, to come up with *one* algorithm that you think is suited for this problem and decide to put ‘in production’.

¹<https://developer.spotify.com/web-api/get-audio-features/>

Table 1: Details on the available features (from the Spotify API documentation)

Name	Type	Description
acousticness	float	A confidence measure from 0.0 to 1.0 of whether the track is acoustic. 1.0 represents high confidence the track is acoustic.
danceability	float	Danceability describes how suitable a track is for dancing based on a combination of musical elements including tempo, rhythm stability, beat strength, and overall regularity. A value of 0.0 is least danceable and 1.0 is most danceable.
duration_ms	int	The duration of the track in milliseconds.
energy	float	Energy is a measure from 0.0 to 1.0 and represents a perceptual measure of intensity and activity. Typically, energetic tracks feel fast, loud, and noisy. For example, death metal has high energy, while a Bach prelude scores low on the scale. Perceptual features contributing to this attribute include dynamic range, perceived loudness, timbre, onset rate, and general entropy.
instrumentalness	float	Predicts whether a track contains no vocals. "Ooh" and "aah" sounds are treated as instrumental in this context. Rap or spoken word tracks are clearly "vocal". The closer the instrumentalness value is to 1.0, the greater likelihood the track contains no vocal content. Values above 0.5 are intended to represent instrumental tracks, but confidence is higher as the value approaches 1.0.
key	int	The key the track is in. Integers map to pitches using standard Pitch Class notation. E.g. 0 = C, 1 = C major/D minor, 2 = D, and so on.
liveness	float	Detects the presence of an audience in the recording. Higher liveness values represent an increased probability that the track was performed live. A value above 0.8 provides strong likelihood that the track is live.
loudness	float	The overall loudness of a track in decibels (dB). Loudness values are averaged across the entire track and are useful for comparing relative loudness of tracks. Loudness is the quality of a sound that is the primary psychological correlate of physical strength (amplitude). Values typical range between -60 and 0 db.
mode	string	Mode indicates the modality (major or minor) of a track, the type of scale from which its melodic content is derived.
speechiness	float	Speechiness detects the presence of spoken words in a track. The more exclusively speech-like the recording (e.g. talk show, audio book, poetry), the closer to 1.0 the attribute value. Values above 0.66 describe tracks that are probably made entirely of spoken words. Values between 0.33 and 0.66 describe tracks that may contain both music and speech, either in sections or layered, including such cases as rap music. Values below 0.33 most likely represent music and other non-speech-like tracks.
tempo	float	The overall estimated tempo of a track in beats per minute (BPM). In musical terminology, tempo is the speed or pace of a given piece and derives directly from the average beat duration.
time_signature	int	An estimated overall time signature of a track. The time signature (meter) is a notational convention to specify how many beats are in each bar (or measure).
valence	float	A measure from 0.0 to 1.0 describing the musical positiveness conveyed by a track. Tracks with high valence sound more positive (e.g. happy, cheerful, euphoric), while tracks with low valence sound more negative (e.g. sad, depressed, angry).

1.1 Data sets

The data set to classify is available as `songs_to_classify.csv`, and the training data is available as `training_data.csv`. The format of the data files is described in `readme.txt`. For details on the extracted features, see Table 1. The data sets are available at <http://www.it.uu.se/edu/course/homepage/sml/project/>.

1.2 Background

The problem of predicting user preferences is today a hot research topic both in academia and industry: you have probably experienced the "You would perhaps also like ..." in online services you use in your everyday life. Within music, a big player has been the Echo Nest, which founded in 2005 as a research spin-off from the MIT Media Lab, and later acquired by Spotify. Their focus is methods for automated understanding of music, and in 2011 they released a popular benchmark dataset 'the million song dataset' (Bertin-Mahieux et al. 2011) which has become popular in the research community (see, for example, Fu et al. 2011; Oord, Dieleman, and Schrauwen 2013) and has similarities to this project. An overview of the scientific field of music recommendation is found in Kaminskas and Ricci (2012).

2 Technical tasks

2.1 Methods to explore

The course has (so far²) covered the five following 'families' of classification methods:

- (i) logistic regression
- (ii) discriminant analysis: LDA, QDA
- (iii) K-nearest neighbor
- (iv) Tree-based methods: classification trees, random forests, bagging
- (v) Boosting

²Neural networks and Gaussian processes, which will be covered later, are also possible to use for classification. You are, of course, welcome to explore these methods as well, in addition to the minimum requirements.

In this project, you decide upon *at least* as many ‘families’ as you are group members, and decide in each ‘family’ *at least* one method to explore. Thus, if you are 4 group members, you explore at least 4 methods from at least 4 different ‘families’.

In addition, you should *also* consider the naive classifier that always predicts LIKE.

2.2 What to do with each method

For *each* methods you decide to explore, you should do the following:

- (a) Implement the method using R. You may write your own code or use packages (the material from the problem solving sessions can be useful).
- (b) Evaluate its performance (using, e.g., cross validation).
- (c) Tune the method to perform well.

Thereafter, you should with a good motivation select which method you decide to use ‘in production’. Inspiration for a ‘good motivation’ can for instance be found in Section 4.5 in James et al. (2013). When you have decided which method to try ‘in production’, run it on the test data (for which you do not have the true labels) and submit your results to <http://www.it.uu.se/edu/course/homepage/sml/project/submit/> to see how well it performs. You submit your results as a string like 010011011, where 0 means DISLIKE and 1 means LIKE. The web site also contains a leader board, where you can see how well you are doing in predicting Andreas’ music taste, compared to the other groups. Each group can submit only one solution per day (only the latest submission will be considered), and the leader board is only updated at midnight every night.

3 Documentation

You should summarize your work by writing a report, which will be reviewed by your coursemates as well as the teaching assistants.

3.1 What to include in your report

The report should include the following:

- (1) A brief introduction to the problem
- (2) A concise description of each of the considered methods, and how they are applied to the problem. Even if you used high-level commands, such as `glm` for logistic regression, you should explain what is ‘under the hood’ of the command! (Please, use your own words. We already know what is written in the book and on Wikipedia, so do not copy-paste from there. Writing these summaries are for your own learning.)
- (3) How the methods were applied to the data (which inputs were used, were the inputs considered qualitative or quantitative, how parameters were tuned, etc), including motivations of the choices made.
- (4) Your evaluation of how well each method performs on the problem.
- (5) Which method you decide to use ‘in production’, and your (good) arguments for your choice!
- (6) How well your method performs ‘in production’ (as obtained from <http://www.it.uu.se/edu/course/homepage/sml/project/submit/>).
- (7) Your conclusions.
- (8) The contribution of each group member (e.g., "Ada implemented KNN, Edith implemented boosting, and Grace was responsible for designing the cross validation to compare the different methods.")
- (9) Appropriate references.
- (10) All code needed to reproduce your reported findings (in an appendix).

The reports are submitted using the student portal.

3.2 How to format your report

Your report should be submitted as a PDF-file following the style used for the prestigious machine learning conference Neural Information Processing Systems (NIPS), which also is the style used for this document. In the NIPS format, your report should be *no longer than 6 pages* (not counting the reference list and code appendix). Except for the page limitation, you should follow the NIPS style closely, including its instructions for figures, tables, citations, etc.

The recommended word processor to use is Latex. You can access the Latex files from the conference webpage <https://nips.cc/Conferences/2016/PaperInformation/StyleFiles>. If you prefer not to install a Latex compiler on your computer, you can use online services such as ShareLatex (<https://www.sharelatex.com/>) or Overleaf (<https://www.overleaf.com/>). In your .tex-file, add the lines

```
\makeatletter
\renewcommand{\@noticestring}{}
\makeatother
```

before `\begin{document}` to suppress the conference-specific footnote.

If you instead prefer using Microsoft Word, OpenOffice, LibreOffice or similar, you may use the older NIPS style files available at <https://nips.cc/Conferences/2015/PaperInformation/StyleFiles>.

When you submit your report for the first time, you should *not* include your own names in the report (since it will be reviewed anonymously by your colleagues)! This is the default status in the Latex files. When you later submit your revised report, however, you should include your names. In Latex, this is achieved by the `final` option, i.e., use `\usepackage[final]{nips_2016}`. In the .docx or .rtf format, you have to do the changes manually.

3.3 The title of your report

The title of your report should read "Group SXXX: Mini Project Report", where SXXX is the submission-id you will receive in a separate e-mail from the teaching assistants.

3.4 How to submit your report

You submit your report using the student portal. (PhD students not registered via the student portal submit per e-mail to andreas.svensson@it.uu.se.)

3.5 Peer review of the reports

Your report will be reviewed by another group. You will also receive the report of another group, which you have to review. As a peer reviewer, you are expected to comment on the following aspects of the report:

- (I) Whether the subset of methods chosen to explore is sufficiently large (methods from at least as many ‘families’ as there is group members, plus the method of always predicting LIKE)
- (II) If all tasks (a-c) from Section 2.2 are made for each method.
- (III) Make an assessment of the technical quality of the proposed solution. Have the considered methods been used in a correct way to address the problem at hand? Are there any flaws in the reasoning and/or motivations used?
- (IV) The report includes everything required from Section 3.1.
- (V) The quality of the language in the report is satisfactory.

The review process is “double blind”, meaning that both the original report and the review report should be anonymous. The review report should be submitted as a simple text file (.txt). Start your review report with the following lines:

```
Peer review of the report from group SYYY
The review is done by group SXXX
```

Here, SXXX is your submission-id, and SYYY is the submission-id of the group whose report you have reviewed. There is no strict page limit for the review report, but a guideline is 1/2–1 page in plain text.

The review report is also submitted via the student portal.

Of course, you should use a polite and constructive language in your review. (Tip: *think about how you would assess your own report before you submit it!*)

After you have submitted your review of the other report, you will get the review on your report from another group, as well as comments from the teaching assistants. Depending on the comments, you might be requested to hand in a revised version of the report.

4 Tip for getting started

Most (if not all) background (methods, theory, R commands, etc.) you need to complete the mini project have been/will be covered at lecture 3, 4, 5 and 6 and problem solving session 3, 4, 5 and 6. However, you only need the material from 4 to start with (a) in Section 2.2.

5 Deadlines & other important dates

Moment	Date
Q&A session (optional)	Monday February 13, 8.15-10.00
Last chance to submit to http://www.it.uu.se/edu/course/homepage/sml/project/submit/	Sunday February 19
Submission of report	Monday February 20
Receive the report of another group to peer review	Tuesday February 21
Submit your peer review of another group's report	Monday February 27
Receive peer review and teaching assistants' comments	Friday March 3 (at the latest)
Hand in revised report (if required)	Friday March 17

All deadlines are one minute before midnight, i.e., 23:59. All tasks in the table above (except for the Q&A session) are mandatory in order to pass the project.

Good luck!

References

- Bertin-Mahieux, Thierry, Daniel P.W. Ellis, Brian Whitman, and Paul Lamere (2011). "The million song dataset". In: *Proceedings of the 12th International Society for Music Information Retrieval Conference (ISMIR)*.
- Fu, Zhouyu, Guojun Lu, Kai Ming Ting, and Dengsheng Zhang (2011). "A survey of audio-based music classification and annotation". In: *IEEE Transactions on Multimedia* 13.2, pp. 303–319.
- James, Gareth, Daniela Witten, Trevor Hastie, and Robert Tibshirani (2013). *An introduction to statistical learning. With applications in R*. Springer.
- Kaminskas, Marius and Francesco Ricci (2012). "Contextual music information retrieval and recommendation: state of the art and challenges". In: *Computer Science Review* 6, pp. 89–119.
- Oord, Aaron van den, Sander Dieleman, and Benjamin Schrauwen (2013). "Deep content-based music recommendation". In: *Advances in Neural Information Processing Systems 26 (NIPS)*.